

Distributed Sensor Hub

Final Report for the Distributed Systems Course 2025/2026

Alex Santini

`alex.santini2@studio.unibo.it`

April 23, 2026

Distributed Sensor Hub is a peer-to-peer platform that collects and disseminates sensor telemetry across multiple nodes, ensuring eventual convergence of replicated state without a central coordinator, and providing a continuously updated global view even under dynamic membership changes and partial failures.

The system combines gossip-based membership, failure detection, and push-pull state synchronization to maintain replica alignment as nodes join, leave, or become temporarily unreachable. It adopts a modular design with pluggable sensor interfaces and configurable topology policies, enabling extensibility and flexible deployment scenarios. Data convergence follows an eventually consistent model based on Last-Writer-Wins (LWW) semantics, formalized as a state-based CRDT to ensure deterministic conflict resolution and monotonic replica evolution.

Within the CAP space, the design explicitly prioritizes availability and partition tolerance over strong consistency, guaranteeing convergence once communication is restored. To sustain operability under realistic distributed conditions, the platform integrates delta-based dissemination with full-state anti-entropy recovery and exposes per-node observability APIs for inspecting topology, membership, and convergence dynamics in real time.

Empirical fault-injection tests, including crash-restart and temporary network partition scenarios, demonstrate deterministic recovery and stable convergence. This combination of CRDT-based convergence, hybrid recovery mechanisms, and built-in observability provides a robust design for decentralized sensor data replication.

Disclaimer

During the development of the project and the preparation of this report, the author used Large Language Models (LLMs) as support tools for selected coding tasks, as well as for brainstorming, outlining, and improving the clarity of the written text.



Figure 1: This is an example image

All AI-generated contributions were critically reviewed, validated, and, when necessary, revised or corrected by the author. The author takes full responsibility for the design and implementation of the system and for the final content of this report.

Quick $\text{\LaTeX}$ suggestions

If you need to cite a reference, you can use the `cite` command, like providing the BibTeX key of some entry in the `references.bib` file, e.g.: [1].

You can find pre-cooked BibTeX entries for most Computer Science papers on DBLP.

If you need to include an image, let $\text{\LaTeX}$ decide where to put it by using the `figure` environment, with a `includegraphics` command inside it. If you need to reference a figure, use the `label` command to assign a label to the figure, and then use the `cref` command to reference it. Put figures in the `figures` folder. A complete example is shown in fig. 1.

Do **not** put any placement constraint on figures, such as `[h]` or `[h!]`. Same considerations apply for other floating elements (e.g. tables, algorithms, listings, etc.).

Do **not** use the `\\` or `\newline` commands to break lines.

Just leave an empty line between two paragraphs to start a new one. The new line will be automatically indented. This is intended: it is the $\text{\LaTeX}$ way to separate paragraphs.

1 Concept

Here you should explain:

- The type of product developed with that project, for example (non-exhaustive):
 - Application (with GUI, be it mobile, web, or desktop)
 - Command-line application (CLI could be used by humans or scripts)
 - Library

- Web-service(s)
- Use case description:
 - *what* is the software doing?
 - *where* are the users?
 - *when* and *how frequently* do they interact with the system?
 - *how* do they *interact* with the system? which *devices* are they using?
 - does the system need to *store* user's **data**? *which?* *where?*
 - most likely, there will be *multiple* **roles**
- Why is distribution needed?
 - geographically distributed environments?
 - computation speedup?
 - resource sharing?
 - fault tolerance?
 - other reasons?

In any case, please elaborate.

2 Requirements Elicitation and Analysis

- The requirements must explain **what** (not how) the software being produced should do.
 - you should not focus on the particular problems, but exclusively on what you want the application to do.
- Requirements must be clearly identified, and possibly numbered
- Requirements are divided into:
 - **Functional**: some functionality the software should provide to the user
 - **Non-functional**: requirements that do not directly concern behavioural aspects, such as consistency, availability, etc.
 - **Implementation**: constrain the entire phase of system realization, for instance by requiring the use of a specific programming language and/or a specific software tool
 - * these constraints should be adequately justified by political / economic / administrative reasons...
 - * ... otherwise, implementation choices should emerge *as a consequence of* design

- If there are domain-specific terms, these should be explained in a glossary
- Each requirement must have its own **acceptance criterion**
 - these will be important for the validation phase

2.1 Relevant Distributed System Features

Motivate which distributed system features are relevant for your project, and which are not.

- transparency
 - does your system need to hide distribution details from users or developers?
 - is it important that failures, location, or replication are invisible?
- fault tolerance, dependability – availability, reliability, integrity, maintainability, safety
 - what happens if a component fails? is uninterrupted service required?
 - is data loss or corruption unacceptable?
 - how quickly must the system recover from faults?
- scalability
 - will the system need to handle increasing numbers of users, requests, or data?
 - is it expected to grow over time?
- security, trust
 - is sensitive data being processed or stored?
 - are there multiple user roles with different permissions?
 - is authentication or authorization required?
- resource sharing
 - do multiple users or components need access to shared resources?
 - is coordination or synchronization needed?
- openness, interoperability, heterogeneity of components
 - Will your system interact with external systems or use components built with different technologies?
 - Is standardization or compatibility important?
- evolvability, maintainability
 - Will the system need to be updated or extended after deployment?
 - Is long-term maintenance a concern?

- performance, concurrency, computation / communication efficiency, bandwidth
 - Are there strict requirements on response time or throughput?
 - Will many operations happen in parallel?
 - Is network usage a concern?
- economy, costs
 - Are there budget constraints for development, deployment, or operation?
 - Is minimizing resource usage important?

3 Design

This chapter explains the strategies used to meet the requirements identified in the analysis. Ideally, the design should be the same, regardless of the technological choices made during the implementation phase.

You can re-arrange the sections as you prefer, but all the sections must be present in the end

Important: try to motivate your design choices in relation to the requirements and features identified in section 2 and section 2.1.

3.1 Architecture

- Which architectural style?
 - why?

3.2 Infrastructure

- are there *infrastructural components* that need to be introduced? *how many*?
 - e.g. *clients, servers, load balancers, caches, databases, message brokers, queues, workers, proxies, firewalls, CDNs, etc.*
- how do components *distribute* over the network? *where*?
 - e.g. do servers / brokers / databases / etc. sit on the same machine? on the same network? on the same datacenter? on the same continent?
- how do components *find* each other?
 - how to *name* components?
 - e.g. DNS, *service discovery, load balancing, etc.*

Component diagrams are welcome here

3.3 Modelling

- which **domain entities** are there?
 - e.g. *users, products, orders, etc.*
- how do *domain entities* **map to** *infrastructural components*?
 - e.g. state of a video game on central server, while inputs/representations on clients
 - e.g. where to store messages in an IM app? for how long?
- which **domain events** are there?
 - e.g. *user registered, product added to cart, order placed, etc.*
- which sorts of **messages** are exchanged?
 - e.g. *commands, events, queries, etc.*
- what information does the **state** of the system comprehend
 - e.g. *users' data, products' data, orders' data, etc.*

Class diagram are welcome here

3.4 Interaction

- how do components *communicate*? *when?* *what?*
- *which* **interaction patterns** do they enact?

Sequence diagrams are welcome here

3.5 Behaviour

- how does *each* component **behave** individually (e.g. in *response* to *events* or *messages*)?
 - some components may be *stateful*, others *stateless*
- which components are in charge of updating the **state** of the system? *when?* *how?*

State diagrams are welcome here

3.6 Data and Consistency Issues

- Is there any data that needs to be stored?
 - *what* data? *where?* *why?*
- how should *persistent data* be **stored**?

- e.g. relations, documents, key-value, graph, etc.
 - why?
- Which components perform queries on the database?
 - *when? which queries? why?*
 - concurrent read? concurrent write? why?
- Is there any data that needs to be shared between components?
 - *why? what data?*

3.7 Fault-Tolerance

- Is there any form of data **replication** / federation / sharing?
 - *why? how* does it work?
- Is there any **heart-beating, timeout, retry mechanism**?
 - *why? among* which components? *how* does it work?
- Is there any form of **error handling**?
 - *what* happens when a component fails? *why? how?*

3.8 Availability

- Is there any **caching** mechanism?
 - *where? why?*
- Is there any form of **load balancing**?
 - *where? why?*
- In case of **network partitioning**, how does the system behave?
 - *why? how?*

3.9 Security

- Is there any form of **authentication**?
 - *where? why?*
- Is there any form of **authorization**?
 - which sort of *access control*?
 - which sorts of users / *roles*? which *access rights*?
- Are **cryptographic schemas** being used?
 - e.g. token verification,
 - e.g. data encryption, etc.

4 Implementation

Please report here all the implementation (technology-dependent) choices you made while implementing your design.

If you run out of time for this project, you may consider leaving some aspect unimplemented, and simply discuss how you would have implemented them. In this case, better would be to discuss unimplemented features in section 9.

- which **network protocols** to use?
 - e.g. UDP, TCP, HTTP, WebSockets, gRPC, XMPP, AMQP, MQTT, etc.
- how should *in-transit data* be **represented**?
 - e.g. JSON, XML, YAML, Protocol Buffers, etc.
- how should *databases* be **queried**?
 - e.g. SQL, NoSQL, etc.
- how should components be *authenticated*?
 - e.g. OAuth, JWT, etc.
- how should components be *authorized*?
 - e.g. RBAC, ABAC, etc.

4.1 Technological details

- any particular *framework* / *technology* being exploited goes here

5 Validation

5.1 Automatic Testing

- how were individual components **unit-tested**?
- how was communication, interaction, and/or integration among components tested?
- how to **end-to-end-test** the system?
 - e.g. production vs. test environment
- for each test specify:
 - rationale of individual tests
 - how were the test automated
 - how to run them
 - which requirement they are testing, if any

recall that *deployment automation* is commonly used to *test* the system in *production-like* environment

recall to test corner cases (crashes, errors, etc.)

5.2 Acceptance test

- did you perform any *manual* testing?
 - what did you test?
 - why wasn't it automatic?

6 Deployment

- should one install your software from scratch, how to do it?
 - provide instructions
 - provide expected outcomes
- what software should be installed on the machines to run your project? which versions?
- should one set environment variables or configuration files?
- if you're using containerization (e.g. Docker), describe how you engineered your deployment scrips (e.g. `docker-compose.yml` files)

7 User Guide

- how to use your software?
 - provide instructions
 - provide expected outcomes
 - provide screenshots if possible

8 Self-evaluation

- An individual section is required for each member of the group
- Each member must self-evaluate their work, listing the strengths and weaknesses of the product
- Each member must describe their role within the group as objectively as possible.

It should be noted that each student is only responsible for their own section.

9 Future works

Discuss possible future works, improvements, extensions, optimizations, etc.

Also discuss here any unimplemented feature, and how you would have implemented it.

References

- [1] D. Adams. *The Hitchhiker's Guide to the Galaxy*. San Val, 1995.